

STATS 202A – FINAL PROJECT TASK 3

NAME: ANOOSHA SAGAR

STUDENT ID: 605028604

SUPPORT VECTOR MACHINES

CONSOLE OUTPUT


```
C:\Users\anoos\Anaconda2\python.exe
E:/UCLA/CourseWork/Fall12017/StatisticsProgramming/Week10/Stats202A_HW9_P2.py

Training Accuracy at Iteration  0 :  0.0
Testing Accuracy at Iteration  0 :  0.934065934066
Training Accuracy at Iteration 100 :  0.992647058824
Testing Accuracy at Iteration 100 :  1.0
Training Accuracy at Iteration 200 :  1.0
Testing Accuracy at Iteration 200 :  1.0
Training Accuracy at Iteration 300 :  1.0
Testing Accuracy at Iteration 300 :  1.0
Training Accuracy at Iteration 400 :  1.0
Testing Accuracy at Iteration 400 :  1.0
Training Accuracy at Iteration 500 :  1.0
Testing Accuracy at Iteration 500 :  1.0
Training Accuracy at Iteration 600 :  1.0
Testing Accuracy at Iteration 600 :  1.0
Training Accuracy at Iteration 700 :  1.0
Testing Accuracy at Iteration 700 :  1.0
Training Accuracy at Iteration 800 :  1.0
Testing Accuracy at Iteration 800 :  1.0
Training Accuracy at Iteration 900 :  1.0
Testing Accuracy at Iteration 900 :  1.0
```

GRAPHICAL REPRESENTATION

```
x, y, z = my_SVM(X_train, Y_train, X_test, Y_test)
plt.xlabel("Number of Iterations --->")
plt.ylabel("Accuracy")
plt.title("Comparison of Testing and Training Accuracy in SVM")
plt.plot(y, 'b--', label="Training Accuracy")
plt.plot(z, label="Testing Accuracy", color='orange')

plt.axis([0,1000,0,1.2])
plt.legend()
plt.grid(True)
plt.show()
```



In machine learning, **support vector machines** are supervised learning models with associated learning algorithms that analyze data used for classification and regression analysis. Given a set of training examples, each marked as belonging to one or the other of two categories, an SVM training algorithm builds a model that assigns new examples to one category or the other, making it a non-probabilistic binary linear classifier.

In addition to performing linear classification, SVMs can efficiently perform a non-linear classification using what is called the kernel trick, implicitly mapping their inputs into high-dimensional feature spaces.

From the above graph, SVM does not overfit and provides excellent results for training and testing accuracy. Also, SVM is resilient to noise.

SVM REACHES PEAK ACCURACY THE QUICKEST

ADAPTIVE BOOSTING (ADABOOST)

CONSOLE OUTPUT

```
C:\Users\anoos\Anaconda2\python.exe
E:/UCLA/CourseWork/Fall12017/StatisticsProgramming/Week10/Stats202A_HW9_P2.py

Training Accuracy at Iteration 0 : 0.904411764706
Testing Accuracy at Iteration 0 : 0.89010989011

Training Accuracy at Iteration 100 : 0.930147058824
Testing Accuracy at Iteration 100 : 0.868131868132

Training Accuracy at Iteration 200 : 0.930147058824
Testing Accuracy at Iteration 200 : 0.857142857143

Training Accuracy at Iteration 300 : 0.930147058824
Testing Accuracy at Iteration 300 : 0.857142857143


Training Accuracy at Iteration 400 : 0.930147058824
Testing Accuracy at Iteration 400 : 0.857142857143

Training Accuracy at Iteration 500 : 0.930147058824
```

```
Testing Accuracy at Iteration 500 : 0.857142857143
Training Accuracy at Iteration 600 : 0.930147058824
Testing Accuracy at Iteration 600 : 0.857142857143
Training Accuracy at Iteration 700 : 0.930147058824
Testing Accuracy at Iteration 700 : 0.857142857143
Training Accuracy at Iteration 800 : 0.930147058824
Testing Accuracy at Iteration 800 : 0.857142857143
Training Accuracy at Iteration 900 : 0.930147058824
Testing Accuracy at Iteration 900 : 0.857142857143
```

GRAPHICAL REPRESENTATION

```
x, y, z = my_Adaboost(X_train, Y_train, X_test, Y_test)
plt.xlabel("Number of Iterations --->")
plt.ylabel("Accuracy")
plt.title("Comparison of Testing and Training Accuracy in Adaboost")
plt.plot(y, 'b--', label="Training Accuracy")
plt.plot(z, label="Testing Accuracy", color='orange')
plt.axis([0,1000,0,1.2])
plt.legend()
plt.grid(True)
plt.show()
```



AdaBoost is a type of "Ensemble Learning" where multiple learners are employed to build a stronger learning algorithm. AdaBoost works by choosing a base algorithm (e.g. decision trees) and iteratively improving it by accounting for the incorrectly classified examples in the training set.

The output of the other learning algorithms ('weak learners') is combined into a weighted sum that represents the final output of the boosted classifier. AdaBoost is adaptive in the sense that subsequent weak learners are tweaked in favor of those instances misclassified by previous classifiers. AdaBoost is sensitive to noisy data

and outliers. In some problems it can be less susceptible to the overfitting problem than other learning algorithms. The individual learners can be weak, but as long as the performance of each one is slightly better than random guessing, the final model can be proven to converge to a strong learner

AdaBoost can handle sparse datasets and therefore, can work with weak classifiers but it shows some variations after it reaches peak classification correctness. Observing the above graph, AdaBoost reaches an effective solution early but becomes sensitive to noise in further iterations.

OUT OF ALL THE TESTED TECHNIQUES, ADABOOST GIVES THE LOWEST ACCURACY.

NEURAL NETWORKS

CONSOLE OUTPUT

```
C:\Users\anoos\Anaconda2\python.exe E:/UCLA/CourseWork/Fall2017/StatisticsProgramming/Final/2_layer_nn.py
Training Accuracy at Iteration  0 :  0.481481481481
Testing Accuracy at Iteration  0 :  0.544444444444
Training Accuracy at Iteration 100 :  0.914814814815
Testing Accuracy at Iteration 100 :  0.922222222222
Training Accuracy at Iteration 200 :  0.962962962963
Testing Accuracy at Iteration 200 :  0.944444444444
Training Accuracy at Iteration 300 :  0.962962962963
Testing Accuracy at Iteration 300 :  0.944444444444
Training Accuracy at Iteration 400 :  0.974074074074
Testing Accuracy at Iteration 400 :  0.966666666667
Training Accuracy at Iteration 500 :  0.981481481481
Testing Accuracy at Iteration 500 :  0.966666666667
Training Accuracy at Iteration 600 :  0.988888888889
Testing Accuracy at Iteration 600 :  0.966666666667
Training Accuracy at Iteration 700 :  0.988888888889
Testing Accuracy at Iteration 700 :  0.966666666667
Training Accuracy at Iteration 800 :  0.988888888889
Testing Accuracy at Iteration 800 :  0.966666666667
Training Accuracy at Iteration 900 :  0.981481481481
Testing Accuracy at Iteration 900 :  0.966666666667
```


GRAPHICAL REPRESENTATION

```
X_train, Y_train, X_test, Y_test = prepare_data()
alpha,beta,acc_train,acc_test=my_NN(X_train,Y_train,X_test,Y_test,num_hidden=50,num_ite
rations=1000,learning_rate=1e-2)
plt.axis([0,1000,0,1.2])
plt.xlabel("Number of Iterations ---->")
plt.ylabel("Accuracy")
plt.title("Comparison of Testing and Training Accuracy in a 2 layer Neural Network")
```

```
plt.plot(acc_train, 'b--', label="Training Accuracy")
plt.plot(acc_test, label="Testing Accuracy", color='orange')

plt.legend()
plt.grid(True)
plt.show()
```

Comparison of Testing and Training Accuracy in a 2 layer Neural Network



Neural networks process information in a similar way the human brain does. The network is composed of many highly interconnected processing elements(neurons) working in parallel to solve a specific problem. Neural networks learn by example. They cannot be programmed to perform a specific task. The examples must be selected carefully otherwise useful time is wasted or even worse the network might be functioning incorrectly. The disadvantage is that because the network finds out how to solve the problem by itself, its operation can be unpredictable.

Neural network is a non-linear classifier, and the hidden layers introduce complexity. Neural networks are also heavily parametric.

ALTHOUGH NEURAL NETWORKS TAKE A LONG TIME TO CONVERGE, THEY PROVIDE GOOD ACCURACY.

POINT TO NOTE: IN THE CASE OF RELU IN NEURAL NETWORKS, IT WAS OBSERVED THAT THE PERFORMANCE OF THE NETWORK ALSO DEPENDS ON THE INITIALIZATION OF THE WEIGHTS. FOR DIFFERENT KINDS OF INITIALIZATION, DIFFERENT RESULTS WERE OBSERVED. THIS BEHAVIOR WAS OBSERVED IN TENSORFLOW AND PYTHON ALSO.

NEURAL NETWORKS USING TENSORFLOW

OVERVIEW

TensorFlow™ is an open source software library for numerical computation using data flow graphs. Nodes in the graph represent mathematical operations, while the graph edges represent the multidimensional data arrays (tensors) communicated between them. The flexible architecture allows you to deploy computation to one or more

CPUs or GPUs in a desktop, server, or mobile device with a single API. TensorFlow was originally developed by researchers and engineers working on the Google Brain Team within Google's Machine Intelligence research organization for the purposes of conducting machine learning and deep neural networks research, but the system is general enough to be applicable in a wide variety of other domains as well.

GENERAL PRINCIPAL FOR BUILDING A NEURAL NETWORK

To build a neural network with Tensorflow, you can follow the steps below:


1. Import the modules you need. 2)
2. Define an `add_layer` function to construct layers.
3. Define the variables and initialize them.
4. Create a session to perform the operation.
5. Build the NN, and send data with placeholders and `feed_dict`
6. Train and test the NN, and observe the results with Tensorboard.

BUILDING A TWO LAYER NEURAL NETWORK WITH RELU AND SOFTMAX AS ACTIVATION FUNCTIONS

PROBLEM STATEMENT

```
Train a two layer neural network to classify the MNIST dataset ##  
## Use Relu as the activation function for the first layer. Use Softmax as the  
activation function for the second layer##  
##  $z = \text{Relu}(x * W_1 + b_1)$  ##  
##  $y = \text{Softmax}(z * W_2 + b_2)$  ##  
# Use cross-entropy as the loss function#
```

DIAGRAM OF THE NEURAL NETWORK



OUTPUT AT EACH LAYER

LAYER 1: RECTIFIED LINEAR UNITS (RELU)

```
W1 = tf.get_variable('w1', [784, 500],
initializer=tf.random_normal_initializer(stddev=0.3))
b1 = tf.get_variable('b1', [1,], initializer=tf.random_normal_initializer(stddev=0.3))
y1 = tf.nn.relu(tf.matmul(x, W1) + b1)
```

LAYER 2: SOFTMAX

```
W2 = tf.get_variable('w2', [500, 10],
initializer=tf.random_normal_initializer(stddev=0.3))
b2 = tf.get_variable('b2', [1, ],
initializer=tf.random_normal_initializer(stddev=0.3))
y2 = tf.nn.softmax(tf.matmul(y1, W2) + b2)
```

CHOOSING OPTIMIZATION FUNCTIONS

GRADIENT DESCENT OPTIMIZER

Gradient descent is a first-order iterative optimization algorithm for finding the minimum of a function. To find a local minimum of a function using gradient descent, one takes steps proportional to the *negative* of the gradient (or of the approximate gradient) of the function at the current point.

ADAM OPTIMIZER

The `tf.train.AdamOptimizer` uses Kingma and Ba's [Adam algorithm](#) to control the learning rate. Adam offers several advantages over the simple `tf.train.GradientDescentOptimizer`. Foremost is that it uses **moving averages of the parameters** (momentum); this enables Adam to use a larger effective step size, and the algorithm will converge to this step size without fine tuning.

The main down side of the algorithm is that Adam requires more computation to be performed for each parameter in each training step (to maintain the moving averages and variance, and calculate the scaled gradient); and more state to be retained for each parameter (approximately tripling the size of the model to store the average and variance for each parameter).

FOR MY NEURAL NETWORK I CHOSE THE GRADIENT DESCENT OPTIMIZER

TRAINING THE NETWORK USING CROSS ENTROPY

```
cross_entropy = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=y2,
labels=y_))
train_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)

for epoch in range(100):
    for i in range(int(mnist.train.num_examples / 100)):
        batch_xs, batch_ys = mnist.train.next_batch(100)
        sess.run(train_step, feed_dict={x: batch_xs, y_: batch_ys})
    print("Epoch Number", epoch)
```

```
correct = tf.equal(tf.argmax(y, 1), tf.argmax(y_, 1))

accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))
print("Accuracy: ", sess.run(accuracy, feed_dict={x:mnist.test.images,
y_:mnist.test.labels}))
```

RESULTS OBTAINED

After 100 epochs of training, my network was able to achieve an accuracy of 0.9671 or 96.71%

CONDENSED CONSOLE OUTPUT

C:\Users\anoos\AppData\Local\Programs\Python\Python36\python.exe
E:/UCLA/CourseWork/Fall2017/StatisticsProgramming/Final/tensorflow1.py

Extracting /tmp/tensorflow/mnist/input_data/train-images-idx3-ubyte.gz

Extracting /tmp/tensorflow/mnist/input_data/train-labels-idx1-ubyte.gz

Extracting /tmp/tensorflow/mnist/input_data/t10k-images-idx3-ubyte.gz

Extracting /tmp/tensorflow/mnist/input_data/t10k-labels-idx1-ubyte.gz

2017-12-14 10:06:43.623539: I C:\tf_jenkins\home\workspace\rel-win\M\windows\PY\36\tensorflow\core\platform\cpu_feature_guard.cc:137] Your CPU supports instructions that this TensorFlow binary was not compiled to use: AVX AVX2

Epoch Number 0	Accuracy: 0.8655	Epoch Number 17
Accuracy: 0.7319	Epoch Number 9	Accuracy: 0.8691
Epoch Number 1	Accuracy: 0.866	Epoch Number 18
Accuracy: 0.7478	Epoch Number 10	Accuracy: 0.8685
Epoch Number 2	Accuracy: 0.866	Epoch Number 19
Accuracy: 0.8409	Epoch Number 11	Accuracy: 0.8692
Epoch Number 3	Accuracy: 0.8673	Epoch Number 20
Accuracy: 0.8494	Epoch Number 12	Accuracy: 0.8697
Epoch Number 4	Accuracy: 0.8672	Epoch Number 21
Accuracy: 0.8547	Epoch Number 13	Accuracy: 0.8696
Epoch Number 5	Accuracy: 0.8673	Epoch Number 22
Accuracy: 0.8577	Epoch Number 14	Accuracy: 0.8702
Epoch Number 6	Accuracy: 0.8684	Epoch Number 23
Accuracy: 0.8615	Epoch Number 15	Accuracy: 0.8699
Epoch Number 7	Accuracy: 0.8683	Epoch Number 24
Accuracy: 0.863	Epoch Number 16	Accuracy: 0.8696
Epoch Number 8	Accuracy: 0.8677	Epoch Number 25

Accuracy: 0.8703

Epoch Number 26

Accuracy: 0.8706

Epoch Number 27

Accuracy: 0.8708

Epoch Number 28

Accuracy: 0.947

Epoch Number 29

Accuracy: 0.9556

Epoch Number 30

Accuracy: 0.9588

Epoch Number 31

Accuracy: 0.9607

Epoch Number 32

Accuracy: 0.9611

Epoch Number 33

Accuracy: 0.9619

Epoch Number 34

Accuracy: 0.9615

Epoch Number 35

Accuracy: 0.963

Epoch Number 36

Accuracy: 0.9634

Epoch Number 37

Accuracy: 0.9637

Epoch Number 38

Accuracy: 0.9636

Epoch Number 39

Accuracy: 0.9637

Epoch Number 40

Accuracy: 0.9644

Epoch Number 41

Accuracy: 0.9642

Epoch Number 42

Accuracy: 0.9638

Epoch Number 43

Accuracy: 0.9647

Epoch Number 44

Accuracy: 0.964

Epoch Number 45

Accuracy: 0.9642

Epoch Number 46

Accuracy: 0.9644

Epoch Number 47

Accuracy: 0.9648

Epoch Number 48

Accuracy: 0.9645

Epoch Number 49

Accuracy: 0.9647

Epoch Number 50

Accuracy: 0.9643

Epoch Number 51

Accuracy: 0.9645

Epoch Number 52

Accuracy: 0.9642

Epoch Number 53

Accuracy: 0.9652

Epoch Number 54

Accuracy: 0.9654

Epoch Number 55

Accuracy: 0.9656

Epoch Number 56

Accuracy: 0.9662

Epoch Number 57

Accuracy: 0.9657

Epoch Number 58

Accuracy: 0.9661

Epoch Number 59

Accuracy: 0.967

Epoch Number 60

Accuracy: 0.9656

Epoch Number 61

Accuracy: 0.9659

Epoch Number 62

Accuracy: 0.9666

Epoch Number 63

Accuracy: 0.9665

Epoch Number 64

Accuracy: 0.9666

Epoch Number 65

Accuracy: 0.966

Epoch Number 66

Accuracy: 0.9671

Epoch Number 67

Accuracy: 0.9672

Epoch Number 68

Accuracy: 0.967

Epoch Number 69

Accuracy: 0.967

Epoch Number 70

Accuracy: 0.967

Epoch Number 71

Accuracy: 0.9666

Epoch Number 72
Accuracy: 0.9671
Epoch Number 73
Accuracy: 0.9667
Epoch Number 74
Accuracy: 0.9666
Epoch Number 75
Accuracy: 0.9668
Epoch Number 76
Accuracy: 0.9667
Epoch Number 77
Accuracy: 0.9668
Epoch Number 78
Accuracy: 0.9664
Epoch Number 79
Accuracy: 0.9666
Epoch Number 80
Accuracy: 0.9669
Epoch Number 81

Accuracy: 0.9673
Epoch Number 82
Accuracy: 0.9667
Epoch Number 83
Accuracy: 0.9671
Epoch Number 84
Accuracy: 0.9671
Epoch Number 85
Accuracy: 0.9671
Epoch Number 86
Accuracy: 0.9671
Epoch Number 87
Accuracy: 0.9675
Epoch Number 88
Accuracy: 0.967
Epoch Number 89
Accuracy: 0.9669
Epoch Number 90
Accuracy: 0.9668

Epoch Number 91
Accuracy: 0.9672
Epoch Number 92
Accuracy: 0.9675
Epoch Number 93
Accuracy: 0.9672
Epoch Number 94
Accuracy: 0.9673
Epoch Number 95
Accuracy: 0.9675
Epoch Number 96
Accuracy: 0.9673
Epoch Number 97
Accuracy: 0.9674
Epoch Number 98
Accuracy: 0.9674
Epoch Number 99
Accuracy: 0.9671
0.9671