

# Assignment 7

Name: Anoosha Sagar

Student ID: 605028604

## Lasso Regression

### Testing Code

```
> n = 50  
> p = 200  
> s = 10  
> X = matrix(rnorm(n * p), nrow=n)  
> lambda_all = (100 : 1) * 10  
> beta_true = matrix(rep(0, p), nrow = p)  
> beta_true[1:s] = 1 : s  
> Y = X %%% beta_true + rnorm(n)  
> beta_all = Lasso(X, Y, lambda_all)  
> ans = glmnet(X, Y)  
> plot(ans)
```

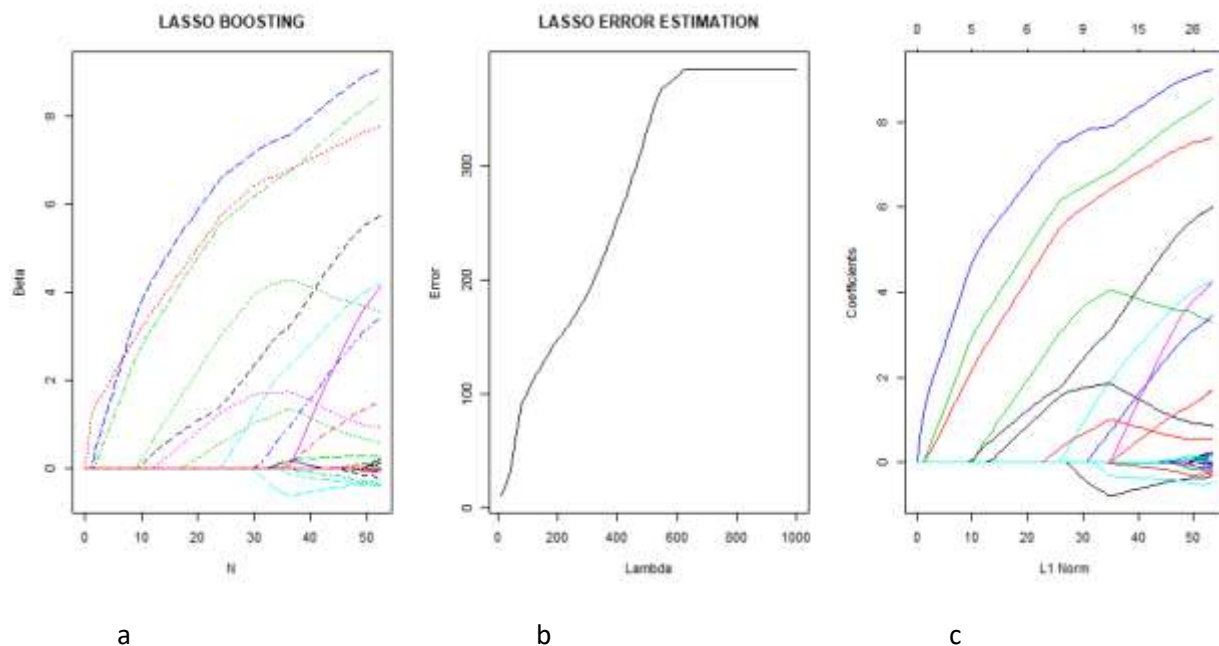


Fig 1: a) Result of Lasso Boosting; b) Lasso Error Estimation; c) Lasso Boosting using glmnet

## Testing with real world data sets

### Linear Regression, Ridge Regression, and Lasso Boosting

*Data Set Used: Swiss Data Set*

*Description:* Standardized fertility measure and socio-economic indicators for each of 47 French-speaking provinces of Switzerland at about 1888.

*Format:*

A data frame with 47 observations on 6 variables, *each* of which is in percent, i.e., in  $[0, 100]$ .

- [,1] Fertility            *Ig*, 'common standardized fertility measure'
- [,2] Agriculture        % of males involved in agriculture as occupation
- [,3] Examination       % draftees receiving highest mark on army examination
- [,4] Education          % education beyond primary school for draftees.
- [,5] Catholic            % 'catholic' (as opposed to 'protestant').
- [,6] Infant.Mortality   live births who live less than 1 year.

All variables but 'Fertility' give proportions of the population.

#### *Snapshot of Data*

```
> swiss <- datasets::swiss
> head(swiss)
```

	Fertility	Agriculture	Examination	Education	Catholic	Infant.Mortality
Courtelay	80.2	17.0	15	12	9.96	22.2
Delemont	83.1	45.1	6	9	84.84	22.2
Franches-Mnt	92.5	39.7	5	5	93.40	20.2
Moutier	85.8	36.5	12	7	33.77	20.3
Neuveville	76.9	43.5	17	15	5.16	20.6
Porrentruy	76.1	35.3	9	7	90.57	26.6

#### Linear Regression

```
> linearRegression(XLin, YLin)
```

\$coefficients	Agriculture	Examination	Education	Catholic	Infant.Mortality
66.9151817	-0.1721140	-0.2580082	-0.8709401	0.1041153	1.0770481

\$standard_error	Agriculture	Examination	Education	Catholic	Infant.Mortality
10.70603759	0.07030392	0.25387820	0.18302860	0.03525785	0.38171965

```
> coef(lm(YLin ~ XLin))
```

(Intercept)	XLinAgriculture	XLinExamination	XLinEducation	XLinCatholic
66.9151817	-0.1721140	-0.2580082	-0.8709401	0.1041153
XLinInfant.Mortality				
1.0770481				

**The values of coefficients obtained using the built in function for linear regression provided by R and the function in my package are similar.**


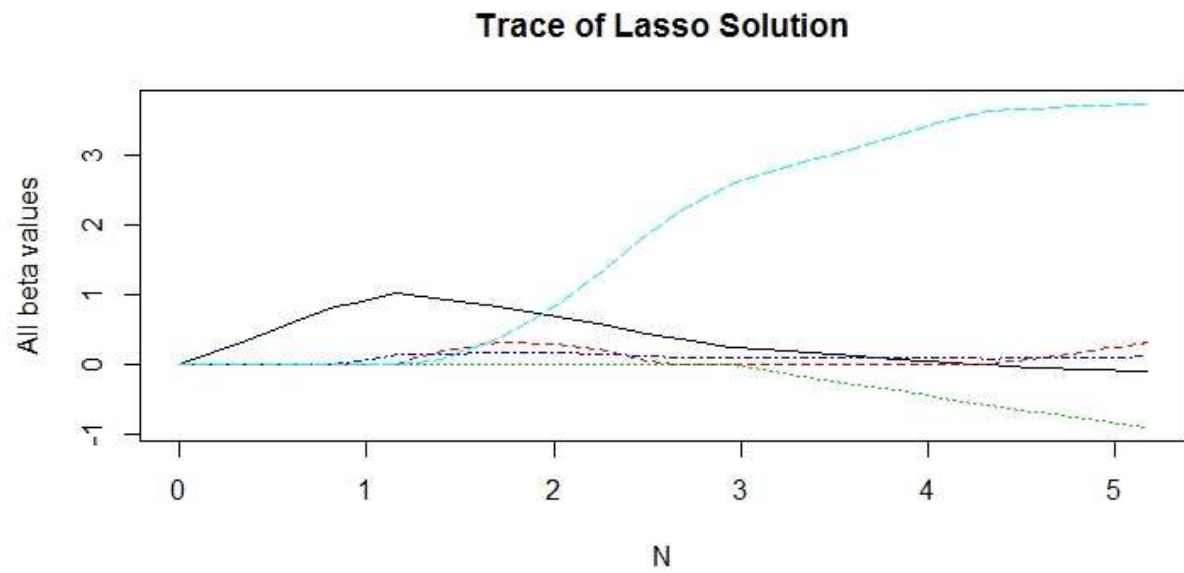
### Ridge Regression

```
> XRidge <- model.matrix(Fertility~., swiss)[-1]
> YRidge <- swiss$Fertility
> lambda <- 10^seq(10, -2, length = 100)
> set.seed(489)
> train <- sample(1: nrow(XRidge), nrow(XRidge)/2)
> test <- (-train)
> yTest <- YRidge[test]
> lambda <- 0.1
> ridgeR <- glmnet(XRidge[train,], YRidge[train], alpha = 0, lambda = lambda)
> ridgeP <- myRidge(x[train,], y[train], lambda)
> ridgeP
[1] 74.64436146 -0.27807670 -0.93900466 -0.35978119 0.06500147 1.37552338
> coef(ridgeR)
6 x 1 sparse Matrix of class "dgCMatrix"
              s0
(Intercept)    73.36350615
Agriculture    -0.26542433
Examination    -0.89519263
Education      -0.36435849
Catholic       0.06570399
Infant.Mortality 1.37394755
```

**The values of coefficients obtained using the built in function for ridge regression provided by R and the function in my package are similar.**

### Lasso Boosting

```
> lambda <- 10^seq(10, -2, length = 100)
> library(Stats202AREgression)
> cv.out <- cv.glmnet(XRidge[train,], YRidge[train], alpha = 0)
> bestlam <- cv.out$lambda.min
> lasso.mod <- glmnet(XRidge[train,], YRidge[train], alpha = 1, lambda = lambda)
> lasso.pred <- predict(lasso.mod, s = bestlam, newx = XRidge[test,])
> mean((lasso.pred-ytest)^2)
[1] 113.7041
> lasso.coef <- predict(lasso.mod, type = 'coefficients', s = bestlam)[1:6,]
> lasso.coef
      (Intercept)      Agriculture      Examination
      57.85476722     -0.06225277     -0.50145205
      Education      Catholic Infant.Mortality
     -0.12425311      0.04456320      1.25231604
> res <- Lasso(XRidge[train,], YRidge[train], lambda)
> plot(lasso.mod)
```



The values of coefficients obtained using the built-in function for lasso boosting provided by R and the function in my package are similar.

As we can observe, the results obtained in all the three scenarios are different. This occurs as linear regression doesn't perform regularization whereas Ridge Regression performs regularization.

Lasso Boosting performs variable selection and regularization to enhance the prediction accuracy.

## Logistic Regression

*Data Set Used: IDRE's Binary Data Set (<https://stats.idre.ucla.edu/stat/data/binary.csv>)*

*Description:* This dataset has a binary response (outcome, dependent) variable called admit. There are three predictor variables: gre, gpa and rank. We will treat the variables gre and gpa as continuous. The variable rank takes on the values 1 through 4. Institutions with a rank of 1 have the highest prestige, while those with a rank of 4 have the lowest.

### *Snapshot of Data*

```
> logData <- read.csv("https://stats.idre.ucla.edu/stat/data/binary.csv")
> head(logData)
  admit gre  gpa rank
1     0 380 3.61    3
2     1 660 3.67    3
3     1 800 4.00    1
4     1 640 3.19    4
5     0 520 2.93    4
6     1 760 3.00    2
```

```
> logData <- read.csv("https://stats.idre.ucla.edu/stat/data/binary.csv")
> X <- as.matrix(logData[, 2:4])
> Y <- as.matrix(logData[, 1])
> X[, 1] <- (X[, 1] - mean(X[, 1])) / sd(X[, 1])
> X[, 2] <- (X[, 2] - mean(X[, 2])) / sd(X[, 2])
> X[, 3] <- (X[, 3] - mean(X[, 3])) / sd(X[, 3])
> logisticRegression(X, Y)

$coefficients
      gre      gpa      rank
0.2233584 0.2510192 -0.4472078
```

```
$standard_error
      gre      gpa      rank
0.1147555 0.1140612 0.1082179
```

```
> print(glm(formula = Y ~ X + 0, family = "binomial"))
```

```
Call: glm(formula = Y ~ X + 0, family = "binomial")
```

```
Coefficients:
      Xgre      Xgpa      Xrank
0.2217    0.2500   -0.4453
```

```
Degrees of Freedom: 400 Total (i.e. Null); 397 Residual
```

```
Null Deviance:          554.5
```

```
Residual Deviance: 519.9    AIC: 525.9
```

The values of coefficients obtained using the built in function for logistic regression provided by R and the function in my package are similar.

## PCA

*Data Set Used: Iris data set*

*Description:* This famous (Fisher's or Anderson's) iris data set gives the measurements in centimeters of the variables sepal length and width and petal length and width, respectively, for 50 flowers from each of 3 species of iris. The species are Iris setosa, versicolor, and virginica.

*Format:*

iris is a data frame with 150 cases (rows) and 5 variables (columns) named Sepal.Length, Sepal.Width, Petal.Length, Petal.Width, and Species.

iris3 gives the same data arranged as a 3-dimensional array of size 50 by 4 by 3, as represented by S-PLUS. The first dimension gives the case number within the species subsample, the second the measurements with names Sepal L., Sepal W., Petal L., and Petal W., and the third the species.

*Snapshot of Data*

```
> iris <- datasets::iris
> head(iris)
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1         3.5          1.4          0.2  setosa
2          4.9         3.0          1.4          0.2  setosa
3          4.7         3.2          1.3          0.2  setosa
4          4.6         3.1          1.5          0.2  setosa
5          5.0         3.6          1.4          0.2  setosa
6          5.4         3.9          1.7          0.4  setosa
> |
```

*Comparison of Results Obtained from PCA function in my package and the built in eigen function*

```
> a <- as.matrix(iris[, 1:4])
> result <- PCA(t(a) %*% a)
> result
$D
[1] 9208.30507 315.45432 11.97804 3.55257

$V
      [,1]      [,2]      [,3]      [,4]
[1,] -0.7511082 -0.2841749 0.50215472 0.3208143
[2,] -0.3800862 -0.5467445 -0.67524332 -0.3172561
[3,] -0.5130089 0.7086646 -0.05916621 -0.4807451
[4,] -0.1679075 0.3436708 -0.53701625 0.7518717

> eigenResult <- eigen(t(a) %*% a)
> eigenResult
eigen() decomposition
$values
[1] 9208.30507 315.45432 11.97804 3.55257

$vectors
      [,1]      [,2]      [,3]      [,4]
[1,] -0.7511082 0.2841749 -0.50215472 0.3208143
[2,] -0.3800862 0.5467445 0.67524332 -0.3172561
[3,] -0.5130089 -0.7086646 0.05916621 -0.4807451
[4,] -0.1679075 -0.3436708 0.53701625 0.7518717
.
```

The eigen values obtained using both the functions obtained is same.